

# Reproducible Research Using Knitr/R

Keith Hughitt

2013/09/03

## Reproducible Research Using Knitr/R

### Overview

This brief tutorial will describe some best practices for using [R](#), knitr, and a couple other related tools for reproducible research.

### Why knitr?

[Knitr](#) is a package for R which allows you to write documents that combine R code with text written in another language such as [LaTeX](#) or [Markdown](#). Knitr then processes the document, running the R code portions and embedding the results in the document, making it easy to generate elegant looking reports in a variety of formats such as Markdown, HTML, and PDF. Unlike the traditional approach where figures are generated separately from the rest of a manuscript, knitr makes it easy to combine all of these elements in such a way so that another scientist can easily reconstruct the entire manuscript from scratch.

### Prerequisites

The main prerequisite for this tutorial is that at least part of your analysis is being performed in R. Some excellent tools exist to achieve a similar effect in other programming languages (e.g. [IPython Notebook](#)), and some of the general practices discussed are relevant regardless of what programming language you are working with, but parts of the tutorial assume that you are working with R.

To get started, let's load up an R console and install and load a few useful packages:

```
install.packages(c("devtools", "knitr", "ggplot2"), dependencies = TRUE)
library(devtools)
```

```
library(knitr)

install_github("knitcitations", "cboettig")
library(knitcitations)
```

Although knitr has support for embedding R code in a number of popular languages, this tutorial focuses on one particular markup language, Markdown, which has a clear and simple syntax, and is supported by many other useful tools such as Github.

## Development Environment

No special development environment is actually needed to work create knitr documents: any text-editor will work. If you are familiar with [RStudio](#), however, there is [built-in support](#) or compiling RMarkdown documents written using the knitr syntax.

## The Basics

The basic structure of an RMarkdown knitr document is to have block of R code, surrounded by Markdown text.

### 1. A simple example

#### Code

Here is some *\*text\** written in [Markdown] (<http://daringfireball.net/projects/markdown/>) followed by a short block of R code.

```
```{r}
print("Hello World")
```
```

Some more text can then follow the code block.

#### Output

Here is some *text* written in [Markdown](#) followed by a short block of R code.

```
print("Hello World")

## [1] "Hello World"
```

Some more text can then follow the code block.

As you can see, the Markdown formatting in the first sentence has been applied, and the R code, in this case a simple print statement, has been evaluated, and the results appending in a quote block following the command.

Although the example is trivial, you can already start to image how this could be useful by including the output of more complex calculations along-side of some explanatory text.

Things get much more interesting, however, when we start to look at how plotting is handled by Knitr.

## 2. A plotting example

Using the same syntax as above, you can also embed plots directly in the output.

For example:

### Code

```
```{r, plot_example, fig.width=8, fig.height=8, fig.dpi=96}
library(ggplot2)
set.seed(1)
x = seq(1, 100)
y = x + rnorm(100, sd=5)

qplot(x, y) + geom_smooth(aes(x, y))
```
```

### Output

```
library(ggplot2)
set.seed(1)
x = seq(1, 100)
y = x + rnorm(100, sd = 5)

qplot(x, y) + geom_smooth(aes(x, y))

## geom_smooth: method="auto" and size of largest group is <1000, so using
## loess. Use 'method = x' to change the smoothing method.
```

For the above example, the figure is generated and saved as a PNG in the **figure** directory. The plot is then referenced in the resulting Markdown, etc. output file so that the image appears directly after the code used to generate the image.




Figure 1: plot of chunk plot\_example

[Chunk options](#) allow you to control various aspects of the code block. In the above example we specified a desired figure width, height, and DPI; other options provide a way to control the verbosity of the code output, skip blocks, etc.’

Notice also in the above example the ‘plot\_example’ string that was added to opening line of the R code block. This is simply an identifier that we can give to that particular block or “chunk” of R code. It is not necessary to give a code chunk an identifier, but it can be helpful for identifying problems down the road, and also results in more descriptive filenames for the plot images.

Finally, whenever parts of your analysis include a random component, such as drawing from the Normal distribution above, it is a good practice to set a known seed: otherwise it is very unlikely that someone else running your code will arrive at the same results.

### 3. Running code written in another language

Another useful feature of knitr is the ability to run [code written in other languages besides R](#).

For example, you can include a [Bash](#) chunk to execute arbitrary shell commands, including other software and scripts. Similar to above, all of the output of those commands will be appended to the knitted report.

#### Code

```
```{r bash_example, engine='bash'}
bowtie2 --version
```
```

#### Output

```
bowtie2 --version

## /usr/bin/bowtie2-align version 2.1.0
## 64-bit
## Built on Europa
## Sat May  4 07:20:58 EDT 2013
## Compiler: gcc version 4.8.0 20130425 (prerelease) (GCC)
## Options: -O3 -m64 -msse2 -funroll-loops -g3
## Sizeof {int, long, long long, void*, size_t, off_t}: {4, 8, 8, 8, 8, 8}
```

### 4. What else can knitr do?

Hopefully the above examples provide an idea of some of the basic functionality of knitr.

Knitr is actually capable of much more, however. For example, you can embed LaTeX commands, images, and movies, and conditionally include child knitr documents.

A full list of features is available in the [knitr documentation](#). An [extensive repository of knitr examples](#) written in Markdown, LaTeX, HTML, etc is also available on Github.

Finally, several other tools exist which can serve to extend the functionality of knitr. One such tool, described below, is knitcitations. Another useful tool for working with the output knitr is [Pandoc](#). Pandoc is a universal document converter, capable of converting between formats such as LaTeX, Markdown, docx, etc. Especially useful is the ability to go from RMarkdown to PDF. For example, to convert this tutorial into a PDF, one could simply do:

```
pandoc README.md -o reproducible_research.pdf
```

Pandoc also has it's own syntax which can be combined with knitr's to produce even nicer results. For example, to create an elegant-looking title page in the Pandoc output, you could include a short comment block at the top of the Markdown file with the following format:

```
% Reproducible Research Using Knitr/R
% [Keith Hughitt] (mailto:khughitt@umd.edu)
% 2013/09/03
```

## Knitcitations

Another useful tool for writing manuscripts directly in Knitr is the [knitcitations package](#). Knitcitations makes it easy to build up and display a bibliography with just the relevant [DOIs](#).

### Code

```
```${r knit_citations}
library(knitcitations)
cleanbib()
citep('10.1126/science.1213847')
citep('10.1093/biostatistics/kxq029')
bibliography()
```
```

### Output

```

library(knitr)
cleanbib()
citep("10.1126/science.1213847")

## [1] "<a href=\"http://dx.doi.org/10.1126/science.1213847\">Peng, 2011</a>"

citep("10.1093/biostatistics/kxq029")

## [1] "<a href=\"http://dx.doi.org/10.1093/biostatistics/kxq029\">Diggle & Zeger, 2010</a>"

bibliography()

##
## - P. J. Diggle, S. L. Zeger, (2010) Editorial. Biostatistics 11 375-375 [10.1093/biostatistics/kxq029]
## - R. D. Peng, (2011) Reproducible Research in Computational Science. Science 334 1226-1227 10.1126/science.1213847

```

Which, when embedded in a Markdown document, looks like:

- P. J. Diggle, S. L. Zeger, (2010) Editorial. *Biostatistics* **11** 375-375 [10.1093/biostatistics/kxq029](http://dx.doi.org/10.1093/biostatistics/kxq029)
- R. D. Peng, (2011) Reproducible Research in Computational Science. *Science* **334** 1226-1227 [10.1126/science.1213847](http://dx.doi.org/10.1126/science.1213847)

## Best Practices for Reproducible Research

Although there has already been [a lot of discussion](#) recently on [reproducible research](#), I thought I would put together a short list of “best practices”.

### 1. Make your data and code available

This is a pretty obvious first step, but people are going to have a difficult time reproducing your results if they are missing either the code or data. Making both of these available in public repositories is already a big first step towards making it easy for other scientists to reproduce your work.

If you are able to share the code used for the analysis openly, then [Github](#) is an excellent free repository that can be used to host your code. Moreover, Github’s built-in support for displaying Markdown makes it especially useful for knitr-based workflows. The present tutorial, for example, is [hosted on Github](#).

## 2. Make all steps of the analysis transparent

The second important factor to make research clear and reproducible is to make all of the different steps of the analysis as transparent as possible.

At a minimum, provenance information should be included for any datasets used in the analysis. Further, if possible, all steps used to prepare the data, however mundane, should also be included directly in the workflow.

Finally, if certain steps in the middle of the analysis involve calling external scripts that you wrote, consider including those scripts inline in a knitr block, or at the very least, include the code for those scripts with your final document.

## 3. Describe the environment used to perform the analysis

In addition to describing any hardware used to collect data, where relevant, the software environment used to perform the analysis should also be described thoroughly. This includes information about the OS, software versions, and library versions.

Fortunately, knitr makes this especially easy since it can collect much of this information for you.

See the **System Information** section below for an example of how this can be performed.

## 4. Be less random


Set a seed for the random number generator in R or whichever language you are using. As described above, the `set.seed()` function can be used in R to achieve this.

## Summary

That's it! Hopefully this tutorial provided you with an idea of some of tools and practices that can be used to make your research as easily reproducible as possible.

With that, I'll leave you with this story of some of the dangers of *non*-reproducible research.





## System Information

```
system("uname -a")
system("python --version")
sessionInfo()

## R version 3.0.1 (2013-05-16)
## Platform: x86_64-unknown-linux-gnu (64-bit)
##
## locale:
##  [1] LC_CTYPE=en_US.UTF-8      LC_NUMERIC=C
##  [3] LC_TIME=en_US.UTF-8      LC_COLLATE=en_US.UTF-8
##  [5] LC_MONETARY=en_US.UTF-8  LC_MESSAGES=en_US.UTF-8
##  [7] LC_PAPER=C               LC_NAME=C
##  [9] LC_ADDRESS=C            LC_TELEPHONE=C
## [11] LC_MEASUREMENT=en_US.UTF-8 LC_IDENTIFICATION=C
##
## attached base packages:
## [1] stats      graphics  grDevices  utils      datasets  methods   base
##
## other attached packages:
## [1] knitr        citations_0.5-0  bibtex_0.3-6  ggplot2_0.9.3.1
```

```
## [4] knitr_1.4.1          colorout_1.0-0      vimcom_0.9-8
## [7] setwidth_1.0-3
##
## loaded via a namespace (and not attached):
## [1] colorspace_1.2-2      dichromat_2.0-0      digest_0.6.3
## [4] evaluate_0.4.7        formatR_0.9          grid_3.0.1
## [7] gtable_0.1.2          httr_0.2             labeling_0.2
## [10] markdown_0.6.3        MASS_7.3-26          munsell_0.4.2
## [13] plyr_1.8              proto_0.3-10         RColorBrewer_1.0-5
## [16] RCurl_1.95-4.1        reshape2_1.2.2       scales_0.2.3
## [19] stringr_0.6.2         tools_3.0.1          XML_3.98-1.1
## [22] xtable_1.7-1
```

## References

- Carl Boettiger, (2013) knitcitations: Citations for knitr markdown files. <https://github.com/cboettig/knitcitations>
- Yihui Xie, (2013) knitr: A general-purpose package for dynamic report generation in R. <http://yihui.name/knitr/>
- Yihui Xie, (2013) Dynamic Documents with {R} and knitr. <http://yihui.name/knitr/>
- Yihui Xie, (2013) knitr: A Comprehensive Tool for Reproducible Research in {R}. <http://www.crcpress.com/product/isbn/9781466561595>>

## Acknowledgements

I would like to thank the authors of knitr and knitcitations, Yihui Xie and Carl Boettiger, for their hard work on these excellent tools, and also for [taking the time](#) to answer questions from pestering knitr newbies.

Keep up the good work!